

Intro to Data Analysis in Python

PyLadies Vancouver Workshop



July 7, 2018

Instructor: Jennifer Walker

Agenda

- Navigating the Python world as a data geek
- Jupyter Lab orientation + quick recap of Python basics
- Working with spreadsheet data
 1. Reading and summarizing CSV files
 2. Basic calculations and graphs
 3. Text data and messy / missing data
 4. Sorting, aggregation, and subsets
- Data visualization: a brief tour of the Python landscape
- Next steps, ideas, and inspiration

Navigating the Python world

... as a data geek

- Python is used in a huge variety of applications

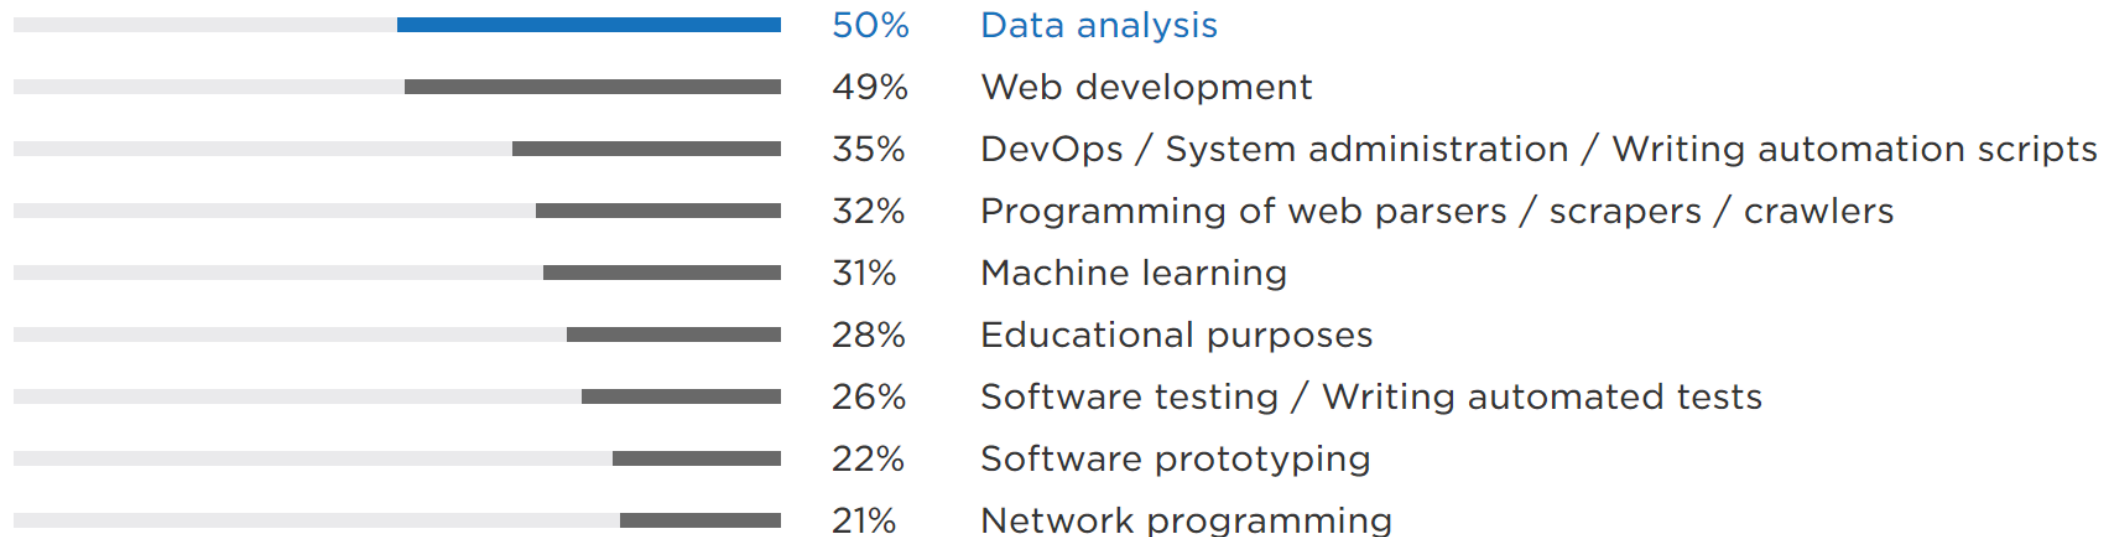


- It has recently become a powerhouse for data analysis

Python Developers Survey 2017

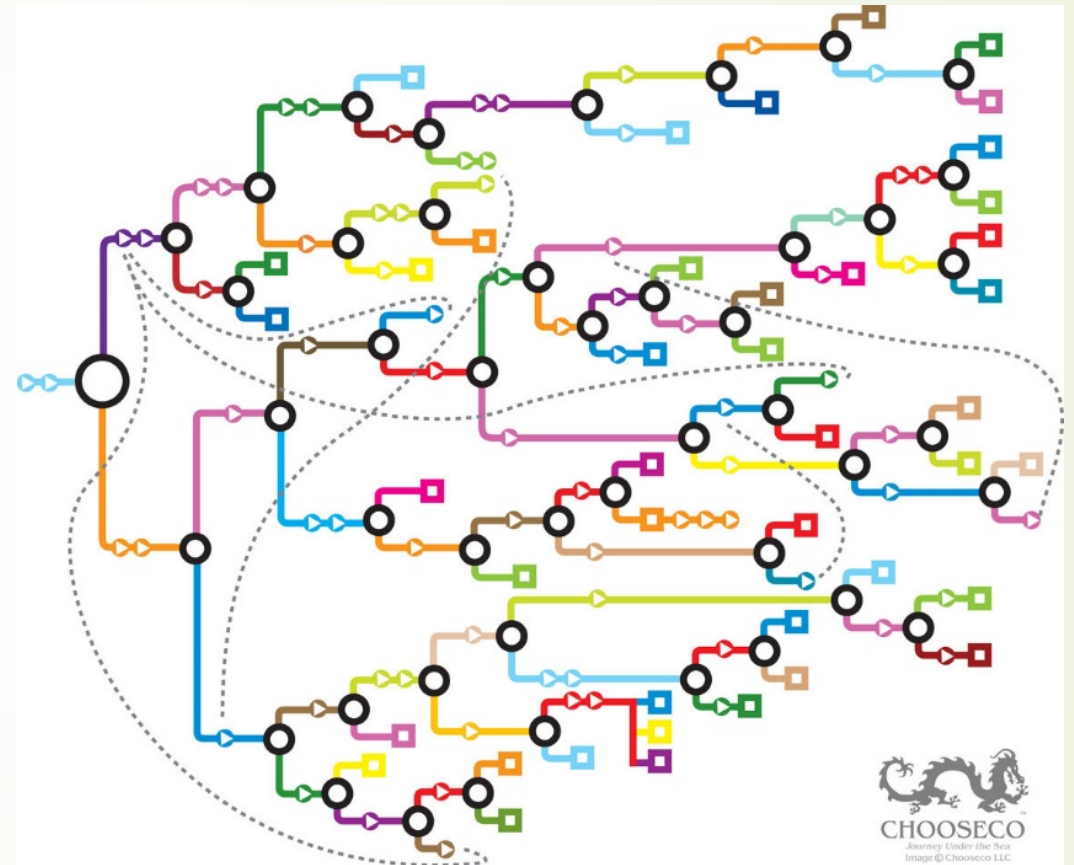

What do you use Python for? (multiple answers)

Combined Python is main Python is secondary




[See all the results](#) ▼

Choose Your Own Adventure




<https://www.atlasobscura.com/articles/cyoa-choose-your-own-adventure-maps>

Python 3 vs. Python 2



Python 3 adoption



Jupyter Lab


- Development environment for working with data
 - Human-centered, interactive coding
- We'll be using Jupyter **notebooks**:
 - Code, graphs, formatted text, equations, etc. in a single document
 - Ideal for exploratory data analysis and for sharing your work with others who are interested in the entire workflow (step by step presentation of your code and results)
 - Uses an IPython kernel to run Python code (IPython = Interactive Python)
 - Many handy features for a much better interactive experience compared to the standard Python console
 - Also supports R, Julia, Perl, and over 100 other languages (and counting!)

Example Jupyter Notebook

<https://www.kaggle.com/arthurtok/generation-unemployed-interactive-plotly-visuals>

Anaconda and Jupyter

What additional technology(s) do you use in addition to Python? (multiple answers)



[See all the results](#) ▼

<https://www.jetbrains.com/research/python-developers-survey-2017>

Python vs. Other Software



Free, open source

Ecosystem of tools developed separately by different teams

Sometimes need to fiddle with settings or implement a few extra steps to get things working

Dynamic, rapidly evolving

- New releases of a tool sometimes introduce drastic changes and incompatibilities with previous versions and other tools



\$\$\$, proprietary, closed source

Self-contained software developed by a single company

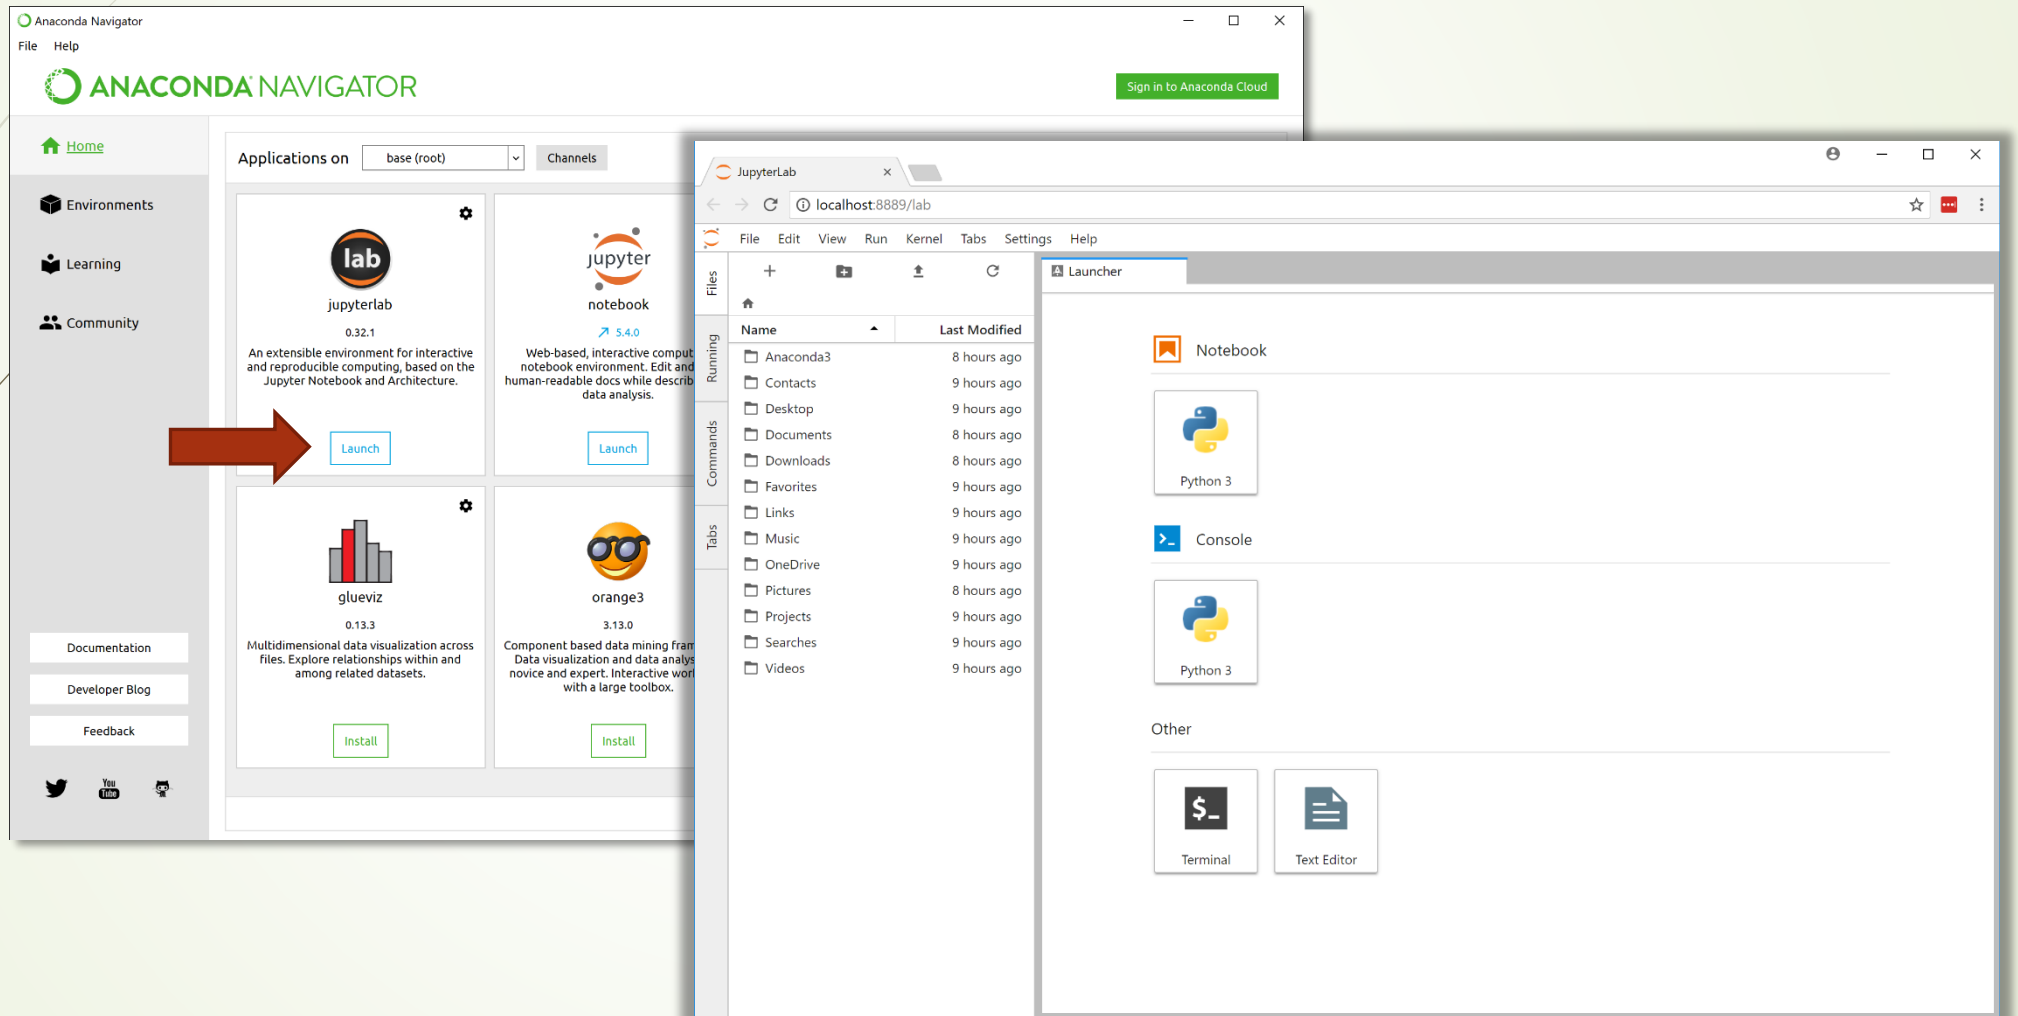
Works "out of the box" on any computer

Gradual and less frequent changes


No need to panic!

- *Google search*
- *Tons of online resources*
- *Huge community of Python users helping each other out*
- *You're not alone!*

Let's Get Coding!






Data Visualization



pandas




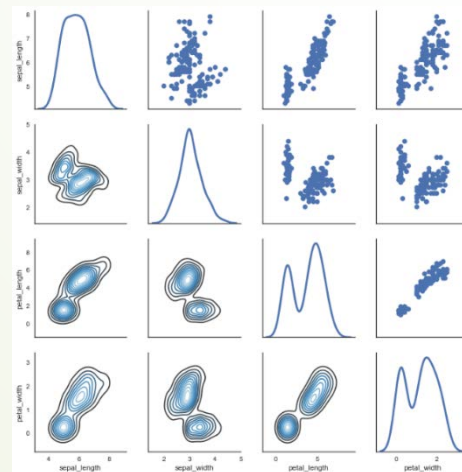


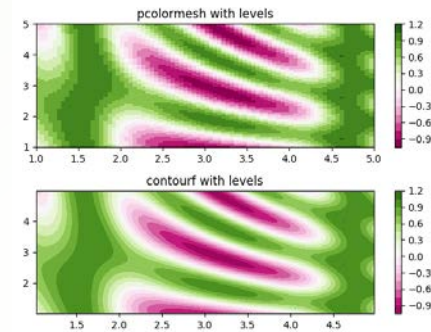



$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$













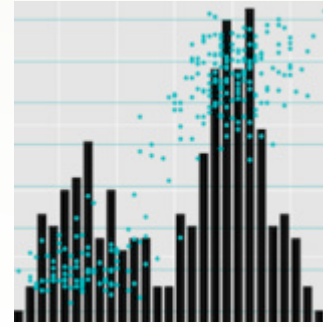


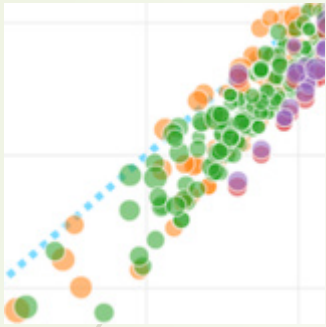
seaborn



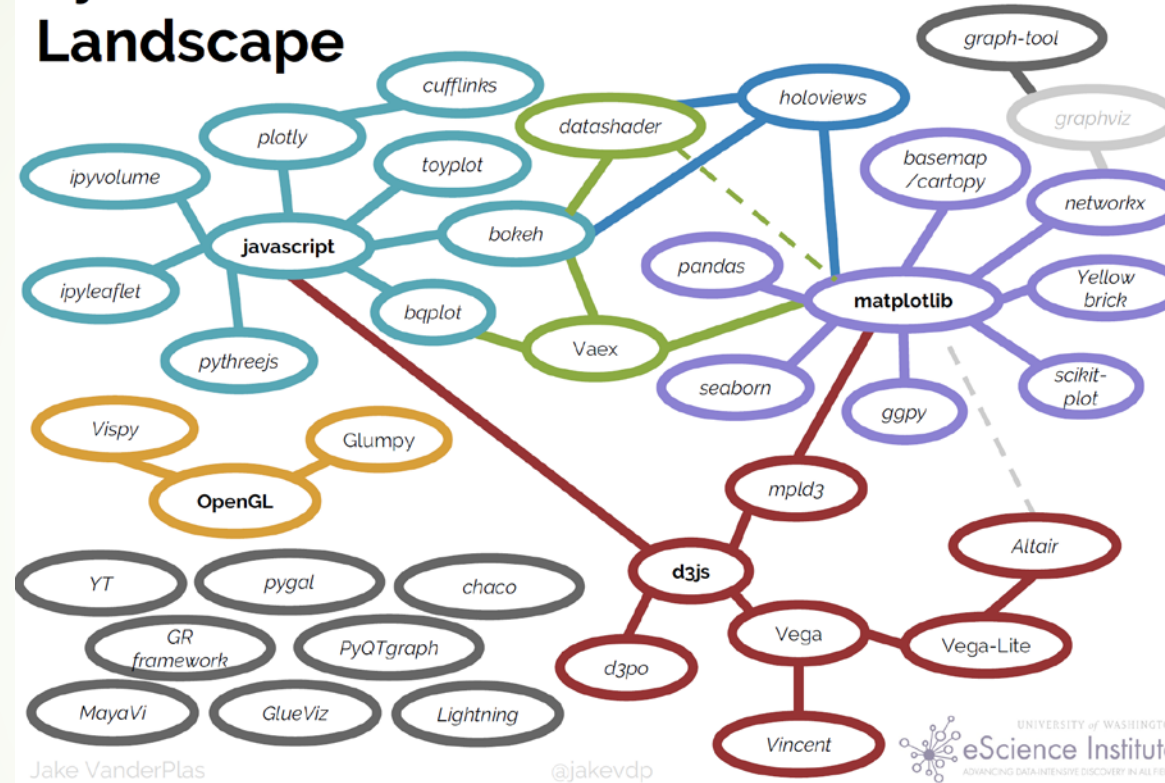
matplotlib & seaborn



plotly



Python's Visualization Landscape



- We've only looked at a tiny fraction of the visualization tools available in Python
- To learn about some of the other amazing visualization libraries and how they fit into the bigger picture, check out this fantastic talk by Jake VanderPlas from PyCon 2017:
<https://www.youtube.com/watch?v=FytuB8nFHPQ>

Visualization Examples & Resources

- pandas: <https://pandas.pydata.org/pandas-docs/stable/visualization.html>
- matplotlib: <https://matplotlib.org/gallery/index.html>
- seaborn: <https://seaborn.pydata.org/examples/index.html>
- Interactive plots:
 - plotly: <https://plot.ly/python/>
 - bokeh: <https://bokeh.pydata.org/en/latest/docs/gallery.html>
- Maps:
 - plotly: <https://plot.ly/python/#maps>
 - cartopy: <https://scitools.org.uk/cartopy/docs/v0.15/gallery.html>
 - folium: <http://folium.readthedocs.io/en/latest/>

Where to go from here?

- Online resources and courses:
 - Data Carpentry: <http://www.datacarpentry.org/python-ecology-lesson/>
 - Data Camp: <https://www.datacamp.com/courses/intro-to-python-for-data-science>
 - Dataquest: <https://www.dataquest.io/>
 - Excellent blog with great tutorials and useful articles: <https://www.dataquest.io/blog/>
 - Kaggle: <https://www.kaggle.com/learn/overview>
 - Many more example Jupyter notebooks and tutorials: <https://www.kaggle.com/kernels>
 - Tons of datasets to play with: <https://www.kaggle.com/datasets>
 - plus Coursera, Udemy, and many others
- Book: Python for Data Analysis, by Wes McKinney. All data and code from the book is at <https://github.com/wesm/pydata-book>

Ideas & Inspiration

- PyData 101 – Jake VanderPlas
 - <https://www.youtube.com/watch?v=DifMYH3iuFw>
- Reproducible Data Analysis in Jupyter – Jake VanderPlas
 - <http://jakevdp.github.io/blog/2017/03/03/reproducible-data-analysis-in-jupyter/>
- Project Jupyter: From Interactive Python to Open Science - Fernando Perez
 - <https://www.youtube.com/watch?v=xuNj5paMuow>
- The Next Generation of Data Products – Hilary Mason
 - <https://www.youtube.com/watch?v=OuRINNSDtIM>

The image features a light beige background with a dark brown vertical bar on the left edge. On the left side, there are several thin, dark brown curved lines that sweep upwards and outwards, resembling stylized grass or reeds. Centered on the slide is the text "Thank You!" in a dark brown, serif font.

Thank You!



Jupyter Lab Orientation

Summary Notes

Getting Started with Jupyter Lab

- Opening Jupyter Lab
 - From Anaconda Navigator: click the “Launch” button under the “lab” icon
 - From the command line (Anaconda Prompt on Windows, Terminal on Mac), run the command: `jupyter lab`
- Jupyter Lab documentation: <http://jupyterlab.readthedocs.io/en/stable/>

Getting Started with Notebooks

- Navigate the “Files” sidebar in Jupyter Lab to change the working directory to your `pydata-intro-workshop` directory for this workshop
- Click the “Files” button to hide / show the “Files” sidebar
- Use the Launcher to create a new Python notebook in the current directory
 - Notice a new file “Untitled.ipynb” appears in the “Files” sidebar
 - The .ipynb extension stands for “IPython notebook” (the original name of the notebook, before it was changed to “Jupyter notebook”)
- Rename your new notebook from “Untitled.ipynb” to “Workshop.ipynb”
 - If the name hasn’t updated in the “Files” sidebar, click the Refresh button

Getting Started with Notebooks

- Notebook auto-saves periodically (like a Google document)
 - Can also manually save at any time with Ctrl-S (Cmd-S on Mac) or by clicking the Save icon on the top toolbar
 - Next time you open up Jupyter Lab, you can return to your "Workshop.ipynb" notebook by double-clicking it in the "Files" sidebar
- A notebook is a collection of cells, where you can run small chunks of code one at a time
 - Our new notebook starts off with a single, empty cell
- Enter some code into the cell (e.g. $2 + 2$) and press Shift-Enter
 - The code runs, displays the output, and advances to the next cell (creating a new cell below if none exists)

Getting Started with Notebooks

- Create some more cells and run some more arithmetic operations
- Notice to the left of the first cell, it says “In[1]” and “Out[1]”, and so on for the second, third, etc. cells
 - The number inside “In[]” and “Out[]” increases by 1 each time you run a cell
 - Try running your last cell again and notice how the number changes
 - Go back to an earlier cell and run it again and notice how the number changes
 - You don’t generally need to worry about these numbers, but they can give you a heads up if you’ve been running your cells out of order

Notebook Structure

- Notebooks can be nonlinear
 - Unlike a script, where the code in a file is executed a single time, from top to bottom, the cells in a notebook can be executed out of order and multiple times each
 - You can also move cells around a notebook by dragging and dropping them
 - You can insert a new cell below a selected cell by clicking the + icon in the top menu
 - Executing code cells out of order can cause unexpected or confusing behaviour
 - This is something to be mindful of when writing notebooks
 - Following some best practices can help you avoid problems

Notebook Kernel

- From the top menu, select Kernel ➔ Restart Kernel and Clear All Outputs, and see what happens in your notebook
- Then from the top menu, select Run ➔ Run All Cells, and see what happens
- The above two steps can be consolidated into a single step: Kernel ➔ Restart Kernel and Run All Cells
 - As you're writing your notebook, periodically running the above command can help ensure that bugs haven't crept in due to running cells out of order
- Check out the items in the Kernel menu and the Run menu for more handy stuff, for example running a selection of several cells, or running all cells above or below a selected cell

Notebook Kernel

- You can see what kernel sessions are currently running by selecting the “Running” sidebar
 - When you close a notebook’s tab in Jupyter Lab, its kernel keeps running in the background
 - To shut it down, you can select “Shutdown” next to the notebook in the “Running” sidebar, or you can quit Jupyter Lab (see later slide)

Formatted Text: Markdown Cells

- Select an empty cell and from the top menu bar, click where it says "Code", and from the dropdown menu select "Markdown"
 - This changes the cell from a Code cell to a Markdown cell, where you can use the Markdown language to quickly format text, insert images, show equations, and more:
 - <https://www.markdownguide.org/getting-started>
 - <https://github.com/adam-p/markdown-here/wiki/Markdown-Cheatsheet>

Formatted Text: Markdown Cells

- In your Markdown cell, type some lines of text such as the example below, then press Shift-Enter to render the cell into formatted text

```
# Python is Awesome

## And so is Jupyter

Here is some regular text. Here are some words in italic and
some words in bold. The area of a circle is  $A = \pi r^2$ .

- Bullet item 1
  - Sub-bullet 1a
- Bullet item 2
- Bullet item 3
```

- To edit a Markdown cell, you need to double click inside it (for a code cell, you only need to single click inside it)
 - Add or remove some lines, and make any other edits you like, and then press Shift-Enter again to display the changes

Splitting and Merging Cells

- A cell can be split into two separate cells
 - Double click inside your Markdown cell to enter edit mode, and put the cursor at the start of one of the lines
 - From the top menu, select Edit ➔ Split Cell
 - Voila! Two separate cells. You can select them and move them around to change the order too.
 - You can also do this with Code cells, when you have multiple lines of code in a cell
- Cells can be merged into a single cell
 - Select two or more Markdown cells, and then from the top menu select Edit ➔ Merge Selected Cells
 - Code cells can be merged the same way

Keyboard Shortcuts

- There are many handy keyboard shortcuts
 - With a cell selected (but not in edit mode)
 - Press Enter to get into edit mode
 - Type A to insert a new cell above or B to insert a new cell below
 - Type M to convert to a Markdown cell or Y to convert to a Code cell
 - Type DD to delete the cell
 - For list of keyboard shortcuts for other commands, check out the “Commands” sidebar

Tabbed and Tiled Views

- You can have multiple notebooks and other files open in Jupyter Lab, and arrange them in different ways
 - From the “Files” sidebar, double-click “Lessons.ipynb” to open it
 - The notebook now shows up as another tab
 - You can keep it in this view, or click on the tab to select it and then start dragging it around the screen to create a tiled view:
 - A blue shaded area appears that changes as you move the cursor around
 - To have a side by side view, drag the cursor until the blue shaded area shows a vertical rectangle on the right. Now the two notebooks are side by side
 - To revert to a single (non-tiled) view, click the notebook’s tab again and drag around until the blue shaded area fills the entire window (excluding the sidebar, if open)

Quitting Jupyter Lab

- To quit Jupyter Lab:
 - Close the Jupyter Lab tab in your browser, and
 - Shutdown the Jupyter Lab server
 - If you launched Jupyter Lab from Anaconda Navigator, close the Anaconda Navigator window (or select File ➔ Quit from the top menu in Anaconda Navigator)
 - If you launched Jupyter Lab from a console (Anaconda Prompt or Terminal), press Ctrl-C (or Cmd-C) at the command line to shutdown the server while keeping the console running, or simply close the console window